Gautam Ajit . Jarvis Consulting

Driven by a passion for transforming data into intelligent solutions, I am a Machine Learning and Data Science professional skilled in Python, SQL, ML, and analytics. I am particularly excited by ML for its ability to uncover patterns, enable prediction, and continuously improve decision-making. Through the Jarvis Machine Learning Engineering program, I built Linux Cluster Monitoring tools, ETL workflows using Pandas and NumPy, and worked on retail store data to improve customer satisfaction and retention, strengthening my data engineering and analytics capabilities. I further explored DLT & ETL pipelines through Azure Databricks, working on stock market and financial transactional data, and recently built an end-to-end pipeline to improve credit scoring and assess how likely is an applicant to default on their loan. I enjoy developing predictive models and leveraging data to solve real-world problems and drive meaningful impact.

Skills

Proficient: Python, Linux/Bash, RDBMS/SQL, Agile/Scrum, Git

Competent: Pandas, NumPy, Databricks, Apache Hive, Tableau

Familiar: LLM APIs, Retrieval-Augmented Generation (RAG), Apache Airflow, Snowflake, R Programming Language

Jarvis Projects

Project source code: https://github.com/jarviscanada/jarvis_data_eng_GautamAjit

Cluster Monitor [GitHub]: Developed a Linux cluster monitoring solution to collect and analyze hardware specifications and server utilization metrics (CPU, memory, and resource usage) for infrastructure capacity planning. Built automated Bash and PostgreSQL scripts to capture and store system data in a Docker-provisioned PostgreSQL database, with Git used for version control. Implemented scheduled monitoring using Cron to collect usage metrics every minute, enabling continuous performance tracking. Designed the MVP on a single node with a scalable architecture capable of supporting multi-node Linux clusters.

Python Data Analytics [GitHub]: Developed a customer analytics proof of concept for a UK-based online retailer to uncover purchasing patterns and support data-driven marketing decisions. Leveraged Python, Pandas, NumPy, and Jupyter Notebook to perform data cleaning, transformation, exploratory data analysis (EDA), and customer segmentation on transactional sales data. Analyzed customer behavior, purchasing trends, and revenue drivers to generate actionable insights for targeted marketing campaigns, customer retention initiatives, and high-value customer acquisition.

Spark on Databricks [GitHub]: Worked on Azure Databricks clusters, ingesting data via the native Lakeflow Connect using a serverless cluster, via Azure Data Factory (ADF) and unity catalog volumes. Implemented the Medallion Architecture to load and transform raw financial transactional data to build a fraud analytics dashboard. Built DLT pipelines to extract stock market data through an API to build a dashboard that visualizes key stock market metrics like Daily price movement, Price trend over 7, 30, and 90 days, Daily percentage change, Trading volume trends & Company-level performance comparisons

Credit Scoring pipeline [GitHub]: Developed a machine learning risk assessment model for an international consumer finance provider to predict loan applicant default likelihood, specifically targeting individuals with limited or no traditional credit history. Integrated primary application data with external credit bureau records to engineer predictive features and evaluate repayment capacity. Optimized classification performance to enable data-driven lending decisions and expand credit access safely.

Highlighted Projects

FedNLP: An Interpretable NLP system to decode Federal Reserve communications: Developed FedNLP, an interpretable Natural Language Processing (NLP) system to predict Federal Funds Rate movements from Federal Open Market Committee (FOMC) communications. Applied advanced NLP techniques and attention-based neural network architectures to analyze financial text, achieving 74% prediction accuracy. Engineered text processing and modeling pipelines to identify influential terms and phrases driving model predictions, improving transparency and enabling deeper insights into monetary policy signals.

Professional Experiences

Machine Learning Engineer, Jarvis Consulting Group (2026-present): Engineered a retail analytics proof of concept to analyze customer purchasing behavior using Python, Pandas, and NumPy, enabling customer segmentation and actionable marketing insights for revenue growth. Developed a Linux cluster monitoring system to collect and analyze CPU, memory, and system usage metrics, leveraging Bash, PostgreSQL, Docker, and Cron for automated data ingestion and storage. Collaborated in agile environments using Git for version control and delivered insights through Jupyter Notebooks and visual analytics.

User Safety & Risk Operations Agent, Accenture (2026): Evaluated and labeled large-scale user-generated content datasets to ensure accuracy, consistency, and compliance with platform policies, supporting downstream data quality and model training processes. Identified and annotated policy-violating or sensitive content with high precision, contributing to reliable ground-truth datasets for automated systems.

Content Moderator, Teleperformance (2023-2025): Moderated and evaluated social media content to ensure compliance with platform policies, achieving 94% accuracy in classification decisions. Maintained high precision on safety guidelines while processing high-volume user-generated content with low latency. Reviewed ads and web content using workflow tools and ensured adherence to moderation standards.

Machine Learning Engineer, OCAS (2023): Built an ETL sandbox environment using Azure Data Factory and Azure Databricks, optimizing data workflows and reducing storage costs by 57%. Developed a data quality framework leveraging Microsoft Purview and PySpark to improve data transparency and reduce data collection errors.

Education

National Institute of Technology Calicut (2012-2016), Bachelor of Technology, Computer Science & Engineering

The University of Sydney (2019-2020), Master of Data Science, School of Computer Science - GPA: 3.86/4.00

Conestoga College (2021-2022), Ontario College Graduate Certificate, Virtualization and Cloud Computing - GPA: 3.90/4.00

Conestoga College (2022-2023), Ontario College Graduate Certificate, Big Data Solution Architecture - GPA: 3.95/4.00

Miscellaneous

- Databricks Certified Data Engineer Professional
- Astronomer Certification for Apache Airflow Fundamentals
- Table tennis player
- Language (Japanese)